

Computer Vision

Bogdan Alexe

bogdan.alexe@fmi.unibuc.ro

University of Bucharest, 2nd semester, 2025-2026

Course structure

1. Low-level vision: Features and filters

Linear filters, color, texture, edge detection, template matching

2. Mid-level vision: Grouping and fitting

Fitting curves and lines, robust fitting, RANSAC, Hough transform, segmentation

3. Mid-level vision: Multiple views

Local invariant feature and description, epipolar geometry and stereo, object instance recognition

4. High-level vision: Object recognition


Object classification, object detection, part based models, bovw models

5. High-level vision: Video understanding

Object tracking, background subtraction, motion descriptors, optical flow

Gradients & edges

Visualization of CNN Feature Maps



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

First layers → Detect edges and simple textures.

Middle layers → Capture textures and object parts.

Final layers → Recognize entire objects.

Images can naturally be decomposed in:

pixel → edge → texton → superpixel → part → object

1.Pixel

- the smallest unit of an image, a single point with a specific color (or intensity in grayscale).
- each pixel doesn't carry much information, it's only when grouping pixels together meaningful structures emerge.

2.Edge

- an edge is a boundary where there is a sharp change in color or intensity.
- by detecting edges, we start to see outlines and boundaries of shapes, the first step in for higher-level features.

3.Texton

- a “texton” is a small, fundamental visual pattern (like a local texture element) which often repeats in a region. For example, the basic brick shape in a brick wall, or a small patch of repeated grain in wood texture.
- textons capture local texture statistic, recognizing them help moving beyond just raw edges to textures/patterns.

4.Superpixel

- a cluster of pixels (often hundreds or thousands) that share similar color or texture. Instead of looking at each individual pixel, you group them into a larger, more meaningful “patch.”
- superpixels simplify the image representation. Each superpixel represents a region that's uniform in some way (color, texture, brightness), which makes further processing more efficient and structured.

5.Part

- a significant portion of an object. For instance, if you have a face, “parts” would be the eyes, nose, or mouth; if you have a car, “parts” might be the wheels, doors, or windshield.
- once you group superpixels into larger, functionally relevant segments, these start to represent meaningful “parts” of an object. Parts are closer to what we humans intuitively recognize.

6.Object

- the complete, recognizable entity in the scene (e.g., a person, a car, a tree).
- combining all the parts (which are themselves made up of superpixels, which in turn come from pixels) yields the full, identifiable object.

Images can naturally be decomposed in:

pixel → edge → texton → superpixel → part → object

1.Pixel

- the smallest unit of an image, a single point with a specific color (or intensity in grayscale).
- each pixel doesn't carry much information, it's only when grouping pixels together meaningful structures emerge.

2.Edge

- an edge is a boundary where there is a sharp change in color or intensity.
- by detecting edges, we start to see outlines and boundaries of shapes, the first step in for higher-level features.

3.Texton

- a “texton” is a small, fundamental visual pattern (like a local texture element) which often repeats in a region. For example, the basic brick shape in a brick wall, or a small patch of repeated grain in wood texture.
- textons capture local texture statistic, recognizing them help moving beyond just raw edges to textures/patterns.

4.Superpixel

- a cluster of pixels (often hundreds or thousands) that share similar color or texture. Instead of looking at each individual pixel, you group them into a larger, more meaningful “patch.”
- superpixels simplify the image representation. Each superpixel represents a region that's uniform in some way (color, texture, brightness), which makes further processing more efficient and structured.

5.Part


- a significant portion of an object. For instance, if you have a face, “parts” would be the eyes, nose, or mouth; if you have a car, “parts” might be the wheels, doors, or windshield.
- once you group superpixels into larger, functionally relevant segments, these start to represent meaningful “parts” of an object. Parts are closer to what we humans intuitively recognize.

6.Object

- the complete, recognizable entity in the scene (e.g., a person, a car, a tree).
- combining all the parts (which are themselves made up of superpixels, which in turn come from pixels) yields the full, identifiable object.

Why do we care about edges?

- Extract information, recognize objects
- Recover geometry and viewpoint



Lab class 2

Automatic grading of multiple choice tests

- a simpler version of a past project

<https://tinyurl.com/CV-2020-Project1>

UNIVERSITATEA DIN BUCUREȘTI

NR. 20

ADMITERE

FACULTATEA DE
MATEMATICĂ ȘI INFORMATICĂ

PROBĂ DE CONCURS

Domeniul: CTI **730**

Sesiunea: IULIE 2018

Nota pe lucrare: 4,90 (palea 5,57) Nota după contestație: 7,30

TEST GRILĂ

MATEMATICĂ

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5			X	
6			X	
7		X		
8				X
9				X
10	X			
11		X		
12		X		
13			X	
14				X
15				X

INFORMATICĂ ☐

FIZICĂ ☒

Număr întrebare	Răspuns			
	A	B	C	D
1	X			
2			X	
3		X		
4	X			
5		X		
6	X			
7	X			
8		X		
9				X
10			X	
11		X		
12				X
13			X	
14	X			
15			X	

NOTĂ : Se bifează X în căsuța corespunzătoare răspunsului corect.

Lab class 2 – Automatic grading of multiple choice tests

MATEMATICA

Număr întrebare	Răspuns			
	A	B	C	D
1	X			
2				X
3		X		
4			X	
5	X			
6			X	
7	X			
8			X	
9		X		
10	X			
11				X
12			X	
13			X	
14		X		
15				X

NOTĂ: Se bifează X în căsuța corespunzătoare

Edge image for query

MATEMATICA


Număr întrebare	Răspuns			
	A	B	C	D
1				
2				
3				
4				
5				
6				
7				
8				
9				
10				
11				
12				
13				
14				
15				

NOTĂ: Se bifează X în căsuța corespunzătoare

Edge image for template


Lab class 1 – Aligning channels in glass plate images

Your first application in Computer Vision





Lab class 1 – Aligning channels in glass plate images

Your first application in Computer Vision



Lab class 1 – Aligning channels in glass plate images



Lab class 2 – Automatic grading of multiple choice tests

MATEMATICA

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5			X	
6			X	
7		X		
8				X
9				X
10	X			
11		X		
12		X		
13			X	
14				X
15				X

NOTĂ : Se bifează X în căuța corectă

Horizontal lines found with Hough transform

MATEMATICA

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5			X	
6			X	
7		X		
8				X
9				X
10	X			
11		X		
12		X		
13			X	
14				X
15				X


NOTĂ : Se bifează X în căuța corectă

Vertical lines found with Hough transform

What causes an edge?

Reflectance change:
appearance
information, texture


Change in surface
orientation: shape



Depth discontinuity:
object boundary

Cast shadows


Edges/gradients and invariance




Closeup of edges



Closeup of edges



Closeup of edges




Closeup of edges




Characterizing edges

- An edge is a place of rapid change in the image intensity function


image



intensity function
(along horizontal scanline)




first derivative




Derivative of a function

- slope of the tangent line to the graph of the function at that point
- **green** > 0
- **black** $= 0$
- **red** < 0



Characterizing edges

- An edge is a place of rapid change in the image intensity function



Derivatives with convolution

For 2D function, $f(x,y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$


For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$


To implement above as convolution, what would be the associated filter?

Derivatives with convolution

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x+1, y) - f(x, y)}{1}$$





Partial derivatives of an image



$$\frac{\partial f(x, y)}{\partial x}$$

$$\frac{\partial f(x, y)}{\partial y}$$



Which shows changes with respect to x?

(showing filters for correlation)

Slide credit: Kristen Grauman

Sobel filters

There exists other approximations for computing partial derivatives:

- Sobel filters: compute derivatives considering larger neighborhoods (3×3)

-1	0	1
-2	0	2
-1	0	1

Vertical Sobel filter
for computing the partial
derivative

$$\frac{\partial f(x, y)}{\partial x}$$

1	2	1
0	0	0
-1	-2	-1

Horizontal Sobel filter
for computing the partial derivative

$$\frac{\partial f(x, y)}{\partial y}$$

Other filters for finite differences

Prewitt: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

Sobel: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

Roberts: $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

```
img = cv2.imread('simona.jpg',0)
sobelx = cv2.Sobel(img,cv2.CV_64F,1,0,ksize=3)
plt.imshow(sobelx,cmap = 'gray')
plt.show()
```






Image gradient

The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid change in intensity



The **gradient direction** (orientation of edge normal) is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

The **edge strength** is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$







Image gradient

The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid change in intensity



The **gradient direction** (orientation of edge normal) is given by:


$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

The **edge strength** is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

or

$$\|\nabla f\| = \left| \frac{\partial f}{\partial x} \right| + \left| \frac{\partial f}{\partial y} \right|$$



Computing the image gradient

1. compute the partial derivatives:

$$\frac{\partial f(x, y)}{\partial x} \quad \frac{\partial f(x, y)}{\partial y}$$

2. compute the gradient magnitude:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2} \quad \text{or} \quad \|\nabla f\| = \left|\frac{\partial f}{\partial x}\right| + \left|\frac{\partial f}{\partial y}\right|$$


Computing the image gradient

1. compute the partial derivative $\frac{\partial f(x, y)}{\partial x}$


Use the vertical
Sobel filter

$M_x =$


-1	0	1
-2	0	2
-1	0	1




Computing the image gradient



Computing the image gradient



Computing the image gradient



Computing the image gradient





image f

$M_x =$


-1	0	1
-2	0	2
-1	0	1



$\frac{\partial f}{\partial x}$

$M_y =$



-1	-2	-1
0	0	0
1	2	1



$\frac{\partial f}{\partial y}$


$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

$[0, 90^\circ]$
 $[90^\circ, 180^\circ]$
 $[180^\circ, 270^\circ]$
 $[270^\circ, 360^\circ]$



From gradients to edges


Gradients



Gradients



Gradients



Sobel filter

Roberts filter

Prewitt filter



Edges

Edges

Edges

Edge detection using 1st derivative

Using first-derivative (gradient-based) methods

- these methods compute the gradient of the image = a vector pointing in the direction of the greatest rate of intensity change.
- common filters: Sobel, Prewitt, and Roberts operators.
- because first derivatives respond strongly to noise, it is common to include smoothing, such as applying a Gaussian blur before taking the gradient (e.g., in Canny edge detection).
- steps typically involve:
 - **compute partial derivatives** in the x and y directions (for example, using a Sobel operator).
 - **calculate the gradient magnitude** from those partial derivatives. A large magnitude indicates a strong edge.
 - **threshold the gradient magnitude** to decide whether an edge is present.

Edge detection using 2nd derivative

Using second-derivative methods

- these methods rely on the second derivative, which measures how quickly the first derivative (the gradient) changes.
- a classic second-derivative operator is the Laplacian:

$$\nabla^2 I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$

- since the second derivative amplifies noise, a smoothing step is almost always applied before or as part of taking the Laplacian. This leads to the **Laplacian of Gaussian (LoG)** filter or the closely related **Difference of Gaussians (DoG)**.
- in a second-derivative approach, an edge often corresponds to **zero-crossings** in the filtered image: where the Laplacian output changes sign (changes in sign +- or -+ in a 4 or 8 connected neighborhood).

Edge detection using 1st and 2nd derivative



Sobel



Roberts






Prewitt





Laplacian

Intensity profile





With a little Gaussian noise



Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal




Where is the edge?

Effects of noise

- Difference filters respond strongly to noise
 - Image noise results in pixels that look very different from their neighbors
 - Generally, the larger the noise the stronger the response
- What can we do about it?


Solution: smooth first



Where is the edge?

Look for peaks in $\frac{\partial}{\partial x}(h \star f)$


Solution: smooth first



Where is the edge?

Look for peaks in $\frac{\partial}{\partial x}(h \star f)$

Solution: smooth first



Where is the edge?

Look for peaks in $\frac{\partial}{\partial x}(h \star f)$ = zero's of the $\frac{\partial^2}{\partial x^2}(h \star f)$
function


Derivative theorem of convolution

$$\frac{\partial}{\partial x}(h \star f) = \left(\frac{\partial}{\partial x}h\right) \star f$$


Differentiation property of convolution.

Sigma = 50


f



$\frac{\partial}{\partial x}h$




$\left(\frac{\partial}{\partial x}h\right) \star f$



1D




d_x



h

$=$




$\frac{\partial h}{\partial x}$


Laplacian of a Gaussian – 1D

Consider $\frac{\partial^2}{\partial x^2}(h \star f)$


f



$\frac{\partial^2}{\partial x^2}h$




$(\frac{\partial^2}{\partial x^2}h) \star f$





Where is the edge?

Zero-crossing of the function $(\frac{\partial^2}{\partial x^2}h) \star f$


2D: Images are noisy




2D: Images are noisy




Solution: smooth first




Smooth first, then derivative




Smooth first, then derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right) *$$



Smooth first, then derivative




Smooth first, then derivative




Smooth first, then derivative



Smooth first, then derivative



Smooth first, then derivative



Smooth first, then derivative



$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$

The diagram illustrates the convolution process. On the left, a 1D kernel $\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$ is convolved with a 2D kernel $\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$. The result is a 2D grid where the top row contains the values -1 , 0 , and 1 , and the bottom two rows are empty. A pink box highlights the top row, and a black arrow points to the right. On the right, the result is scaled by $\frac{1}{2}$, yielding a 2D grid with the top row containing 2 , 0 , and -2 , and the bottom two rows empty.

2	0	-2
4	0	-4

Smooth first, then derivative


$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$



$$\frac{1}{2} \times \begin{bmatrix} 2 & 0 & -2 \\ 4 & 0 & -4 \\ 2 & & \end{bmatrix}$$


Smooth first, then derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$





Smooth first, then derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$



Sobel filter: smooth & derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$



1	0	-1
2	0	-2
1	0	-1

Vertical Sobel filter
for computing the partial
derivative

$$\frac{\partial f(x, y)}{\partial x}$$

Derivative of Gaussian filters - 2D

$$d_x \otimes (h \otimes I) = (d_x \otimes h) \otimes I$$

derivation
filter

Gaussian
filter


image

0	0	0
-1	0	1
0	0	0






0.0030	0.0133	0.0219	0.0133	0.0030
0.0133	0.0596	0.0983	0.0596	0.0133
0.0219	0.0983	0.1621	0.0983	0.0219
0.0133	0.0596	0.0983	0.0596	0.0133
0.0030	0.0133	0.0219	0.0133	0.0030

0	0	0
-1	0	1
0	0	0



=



Derivative of Gaussian filters - 2D


$$d_y \otimes (h \otimes I) = (d_y \otimes h) \otimes I$$

0	-1	0
0	0	0
0	1	0






0.0030	0.0133	0.0219	0.0133	0.0030
0.0133	0.0596	0.0983	0.0596	0.0133
0.0219	0.0983	0.1621	0.0983	0.0219
0.0133	0.0596	0.0983	0.0596	0.0133
0.0030	0.0133	0.0219	0.0133	0.0030


0	-1	0
0	0	0
0	1	0





=




Derivative of Gaussian filters - 2D



x - direction




y - direction




Derivative of Gaussian filters - 2D

0	0	0
-1	0	1
0	0	0

direction x



=




$\frac{\partial h}{\partial x}$


$\frac{\partial^2 h}{\partial x^2}$

0	-1	0
0	0	0
0	1	0

direction y



=




$\frac{\partial h}{\partial y}$


$\frac{\partial^2 h}{\partial y^2}$

Laplacian of Gaussian


$$\nabla^2 h = \frac{\partial^2 h}{\partial x^2} + \frac{\partial^2 h}{\partial y^2}$$



+




=



Example of a
Laplacian filter


0	1	0
1	-4	1
0	1	0




Laplacian of Gaussian

- Laplacian:

$$\Delta f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$




2D edge detection filters



Gaussian


$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$

Laplacian of Gaussian





$$\nabla^2 h_{\sigma}(u, v)$$

- ∇^2 is the Laplacian operator:


$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

Smoothing with a Gaussian


Recall: parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.




...



Effect of σ on derivatives



$\sigma = 1$ pixel



$\sigma = 3$ pixels

The apparent structures differ depending on Gaussian's scale parameter.

Larger values: larger scale edges detected
Smaller values: finer features detected

Mask properties

- Smoothing



- Positive values
- Sum to 1 $\rightarrow$ constant regions same as input
- Amount of smoothing proportional to mask size
- Remove “high-frequency” components; “low-pass” filter

- Derivatives

- Opposite signs used to get high response in regions of high contrast
- Sum to 0 $\rightarrow$ no response in constant regions
- High absolute value at points of high contrast


Another approach: gradient magnitude

- Don't need 2nd derivatives
- Just use magnitude of gradient to compute a gradient map
- Could threshold this gradient map to obtain edges
- Are we done? No!





What we really want: line drawing



Canny Edge Detection

Algorithm:

- Smooth image (only want “real” edges, not noise)
- Calculate gradient direction and magnitude
- Non-maximum suppression perpendicular to edge
- Threshold into strong, weak, no edge
- Connect together components

Smooth image

- Apply a Gaussian filter




Gradient magnitude and direction

- Sobel filter


Horizontal Sobel filter for computing the partial derivative

$$\frac{\partial f(x, y)}{\partial y}$$




Vertical Sobel filter for computing the partial derivative

$$\frac{\partial f(x, y)}{\partial x}$$




Non-maximum suppression


- Want single pixel edges, not thick blurry lines
- Need to check nearby pixels
- See if response is highest




Non-maximum suppression




Non-maximum suppression




Non-maximum suppression




Non-maximum suppression




Non-maximum suppression



Non-maximum suppression



Non-maximum suppression



Non-maximum suppression



Hysteresis Thresholding

After non-maximum suppression, some noise may still remain.

We want to **keep only meaningful edges while discarding noise**.

Idea: Use **two thresholds** instead of one.

Let R be the gradient magnitude.

Three cases:

1. $R > T \rightarrow$ **Strong edge** (definitely an edge)
2. $t < R \leq T \rightarrow$ **Weak edge** (potential edge)
3. $R \leq t \rightarrow$ **Non-edge** (discard)

Why two thresholds?



Hysteresis Thresholding

After non-maximum suppression, some noise may still remain.

We want to **keep only meaningful edges while discarding noise**.


Idea: Use **two thresholds** instead of one.

Let R be the gradient magnitude.

Three cases:

1. $R > T \rightarrow$ **Strong edge** (definitely an edge)
2. $t < R \leq T \rightarrow$ **Weak edge** (potential edge)
3. $R \leq t \rightarrow$ **Non-edge** (discard)

Why two thresholds?



Hysteresis Thresholding

After non-maximum suppression, some noise may still remain.

We want to **keep only meaningful edges while discarding noise**.

Idea: Use **two thresholds** instead of one.

Let R be the gradient magnitude.

Three cases:


1. $R > T \rightarrow$ **Strong edge** (definitely an edge)
2. $t < R \leq T \rightarrow$ **Weak edge** (potential edge)
3. $R \leq t \rightarrow$ **Non-edge** (discard)

Why two thresholds?

A **single threshold** may break real edges.

Weak edges are kept only if connected to strong edges (look in a neighborhood – usually 8 pixels).

This process is called **edge tracking by hysteresis**.




Canny Edge Detection

Algorithm:

- Smooth image (only want “real” edges, not noise)
- Calculate gradient direction and magnitude
- Non-maximum suppression perpendicular to edge
- Threshold into strong, weak, no edge
- Connect together components
- Tunable: Sigma, thresholds

Canny Edge Detection



Canny Edge Detection - demo

<http://bigwww.epfl.ch/demo/ip/demos/edgeDetector/>

Course structure

1. Low-level vision: Features and filters

Linear filters, color, texture, edge detection, template matching

2. Mid-level vision: Grouping and fitting

Fitting curves and lines, robust fitting, RANSAC, Hough transform, segmentation

3. Mid-level vision: Multiple views

Local invariant feature and description, epipolar geometry and stereo, object instance recognition


4. High-level vision: Object recognition

Object classification, object detection, part based models, bovw models

5. High-level vision: Video understanding


Object tracking, background subtraction, motion descriptors, optical flow

Fitting




Fitting

- We have learned how to detect edges
- We would like to form a higher-level, more compact representation of the features in the image by grouping multiple features according to a simple model



Fitting


- Choose a *parametric model* to represent a set of features



simple model: lines



simple model: circles



complex model: car

Lab 2: Line detection - application

PROBĂ DE CONCURS

Domeniul CTI **730**

Sesiunea IULIE 2018

Nota pe lucrare 4,90 (pate 55%) Nota după contestație



TEST GRILĂ

MATEMATICĂ

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5			X	
6			X	
7		X		
8				X
9				X
10	X			
11		X		
12		X		
13			X	
14				X
15				X

INFORMATICĂ ☐

FIZICĂ ☒

Număr întrebare	Răspuns			
	A	B	C	D
1	X			
2			X	
3		X		
4	X			
5		X		
6	X			
7	X			
8		X		
9				X
10			X	
11		X		
12				X
13			X	
14	X			
15			X	

NOTĂ : Se bifează X în căsuța corespunzătoare răspunsului corect.

Lab 2: Line detection - application

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5			X	
6			X	
7		X		
8				X
9				X
10	X			
11		X		
12		X		
13			X	
14				X
15				X

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5			X	
6			X	
7		X		
8				X
9				X
10	X			
11		X		
12		X		
13			X	
14				X
15				X

Lab 2: Line detection - application

MATEMATICA

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5			X	
6			X	
7		X		
8				X
9				X
10	X			
11		X		
12		X		
13			X	
14				X
15				X

MATEMATICA

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5			X	
6			X	
7		X		
8				X
9				X
10	X			
11		X		
12		X		
13			X	
14				X
15				X

Fitting: challenges

- If we know which points belong to the line, how do we find the “optimal” line parameters?
 - Least squares
- What if there are outliers?
 - Robust fitting, RANSAC
- What if there are many lines?
 - Voting methods: RANSAC, Hough transform

Voting schemes

- Let each feature vote for all the models that are compatible with it
 - Hopefully the noise features will not vote consistently for any single model
 - Missing data doesn't matter as long as there are enough features remaining to agree on a good model

Hough transform

- An early type of voting scheme
 - Discretize *parameter space* into bins
 - For each feature point in the image, put a vote in every bin in the parameter space that could have generated this point
 - Find bins that have the most votes







Image space



Hough parameter space

Parameterization of a line




Parameterization of a line in image plane: $y = m * x + b$,

m – slope (defines the angle between the line and Ox axis),

b – intercept (bias to origin $O(0,0)$)

(m,b) are line parameters

Parameterization of a line - examples



$d_1: y = x$, so $m=1$, $b = 0$

$d_2: y = x + 2$, so $m=1$, $b = 2$

$d_3: y = 2x$, so $m=2$, $b = 0$

Parameterization of a line in image plane: $y = m * x + b$,


m – slope (defines the angle between the line and Ox axis),

b – intercept (bias to origin $O(0,0)$)

(m,b) are line parameters

Parameter space representation

- A **line** in the image corresponds to a **point** in Hough space



Hough parameter space





Image space



Hough space

$d_1: y = x$, so $m=1$, $b = 0$

$d_2: y = x + 2$, so $m=1$, $b = 2$

$d_3: y = 2x$, so $m=2$, $b = 0$

$d_1: m=1$, $b = 0 \rightarrow (1,0)$

$d_2: m=1$, $b = 2 \rightarrow (1,2)$


$d_3: m=2$, $b = 0 \rightarrow (2,0)$

A line in the image space corresponds to a point in the Hough space.

Correspondence between image and Hough spaces

A line in the image space corresponds to a point in the Hough space.


What does a **point** (x_0, y_0) in the image space map to in the Hough space?



Correspondence between image and Hough spaces

A line in the image space corresponds to a point in the Hough space.

What does a **point** (x_0, y_0) in the image space map to in the Hough space?




There is an infinity of lines of the form $y = mx + b$ that contain point $(1,1)$:

$y = x$ ($m=1, b=0$); $y = 2x-1$ ($m=2, b=-1$), $y = 1$ ($m=0, b = 1$), $y = -x + 2$ ($m=-1, b = 2$),

Correspondence between image and Hough spaces

A line in the image space corresponds to a point in the Hough space.

A **point** (x_0, y_0) in the image space maps to a line in the Hough space



There is an infinity of lines of the form $y = mx + b$ that contain point $(1,1)$:

$y = x$ ($m=1, b=0$); $y = 2x-1$ ($m=2, b=-1$), $y = 1$ ($m=0, b = 1$), $y = -x + 2$ ($m=-1, b = 2$),

All 4 points $(1,0)$, $(2,-1)$, $(0,-1)$, $(-1,2)$ are on the line $b = -m+1$

Correspondence between image and Hough spaces

A line in the image space corresponds to a point in the Hough space.

A **point** (x_0, y_0) in the image space maps to a line in the Hough space

Parameterization of a line in image plane: $y = m * x + b$,


All lines that contain point (x_0, y_0) can have any slope m and any intercept b , the only condition that should be satisfied is

$$y_0 = m * x_0 + b.$$

$y_0 = m * x_0 + b \Leftrightarrow b = -x_0 * m + y_0$ (parameterization of the line in the Hough space with slope $-x_0$ and intercept y_0)


$$(x_0, y_0) = (1, 1) \Rightarrow b = -m + 1$$

Parameter space representation



Where is the line that contains both (x_0, y_0) and (x_1, y_1) ?

Parameter space representation



Where is the line that contains both (x_0, y_0) and (x_1, y_1) ?

Point (x_0, y_0) maps to the line in the Hough space: $b = -x_0m + y_0$

Point (x_1, y_1) maps to the line in the Hough space: $b = -x_1m + y_1$

The line that contains both (x_0, y_0) and (x_1, y_1) maps to a point in the Hough space. This point is the intersection of lines $b = -x_0m + y_0$ and $b = -x_1m + y_1$